

Lourens Naudé, Trade2Win

# Embracing Events

Stop polling ... react.



RailsWay**Con**

# Bio

- <http://github.com/methodmissing>
  - Ruby contractor
  - General misfit for the current economy
  - Ruby, C#, Java, admin 500 nodes, CSS \* deferred payment
  - Currently building out a consumer forex platform @ Trade2Win  
Slippage matter.
- creator
  - scrooge, mysqlplus\_adapter, mri\_instrumentation
- contributor
  - Rails, mysqlplus, query\_memcached, thinking-sphinx, em-spec

The logo features the text "t2w" in a bold, red, italicized sans-serif font. The letters are slanted upwards to the right. The background is light gray with several thin, curved, white and light orange lines that sweep across the upper left and middle sections, creating a sense of motion or speed.

**t2w**

**RailsWayCon**

# e-vent [i-vent]

something that happens or is regarded as happening; an occurrence, esp. one of some importance.

# A change in state ...

- MySQL server sends results down the wire
  - Connection file descriptor becomes readable
    - Process results
    - Render results
- Ignorance is bliss
  - Never be flooded
  - Only handle events that matter

# Talk#to\_a

- Context switches
- Scheduling
- Programming models
- Frameworks
  - Eventmachine
  - Neverblock
- Operating System
- Lessons from massively multiplayer online games
- Testing
- Q & A ?

# Getting Fibers (coroutines) and Threads

# Scenario we all can relate to ...

- Quiet weekend day, erratic thoughts from the week before comes together
- Sharing space with your better half ... who's being chatty
- You're being interrupted at random
- Let's assume an average "in the zone" startup threshold of 2 minutes per interruption
- Deadlock's not an option ...



# Multithreading

- Let the initial pattern of thought and the interruption be distinct Threads
- The interruption's effective right away
- The context switch is 'expensive' : coding VS whichever chore's been neglected
- Preemptive scheduling == constant interruption

# Fiber / coroutine pool

- Two distinct Fibered contexts
- Able to defer the interruption
- Cheap switching : being thick skinned buys time
- You#resume the interruption if and when there's time to handle it

# Thread

- Scheduling
  - Handled by the VM – timesliced @ 10ms by MRI
  - Executes as soon as CPU resources become available
- Availability
  - MRI 18: Maps to a single native thread
  - MRI 19: 1 x native thread for each Ruby thread - GIL
  - JRuby: 1 x Java thread for each Ruby thread
- Overheads
  - Large initial stack size
  - MRI uses a SIGVTALRM to signal a context switch

# Fiber

- Scheduling
  - No concurrency – scheduled by application code
  - Never executes right away
- Availability
  - MRI 18: Patch by Aman Gupta & Joe Damato
  - MRI 19: Core Fiber class
  - JRuby: Poor Man's Fibers
- Overheads
  - Small initial stack size @ 4k
  - Context switches exclusively in user space – very fast

# Why Fibers matter

- Computes a partial result - generator
- Yields back to the caller
- Saves execution state
- Caller resumes
- Data exchange
  - Accept arguments through `Fiber#resume`
  - Return values through `Fiber.yield`

# Threads as Coroutines

- Through a voodoo combination of
  - Thread priorities
  - Thread.pass
  - Thread.stop
  - Thread#run (yields CPU) or Thread#wakeup (no CPU)
- Poor Man's Fibers piggy backs off a Queue
- MRI 18 Fibers = Threads – preemptive scheduling, much like continuations, but with a smaller stack.
- Performant, but not as lightweight as Fibers

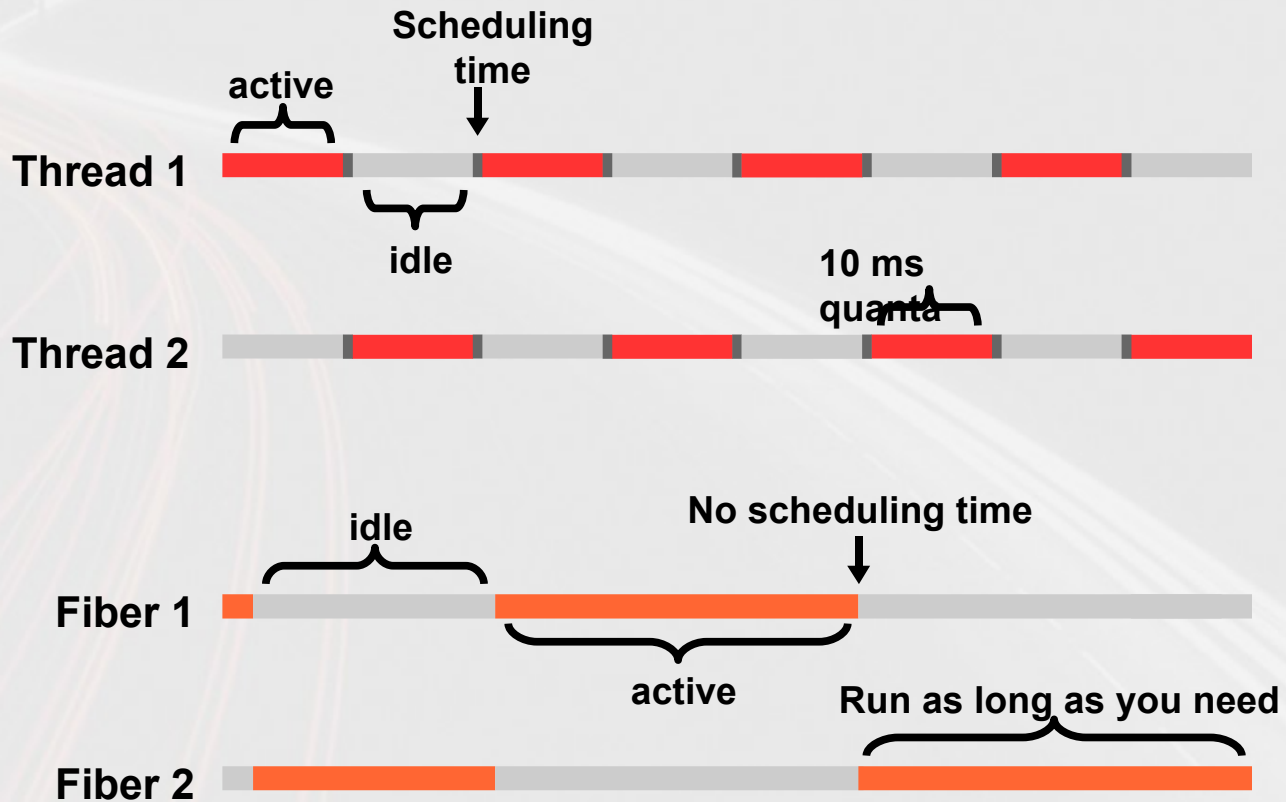
# Cooperative scheduling

- Request / response with evented DB access
  - HTTP request wrapped in a Fiber
  - => mysql->fd : SELECT \* FROM users WHERE id = 1
  - Execution paused
  - <= mysql->fd : #<User id: 1 uname: methodmissing>
  - Execution resumed
  - => mysql->fd: SELECT \* FROM wishlist WHERE user\_id = 1
  - Execution paused
  - <= mysql->fd: #<Wishlist id: 20 user\_id 1>

# Cooperative VS preemptive

- Fibered dispatch, query = \*, result = +, int = &
  - <GET /orders> \* + \*\*++\*\*++\*\*++
  - <POST /checkout> \* + \* + \*\*++
  - <DELETE /sessions> \*+
- Threaded dispatch, interrupted every 10ms
  - <GET /orders> \* & + \*\*++&\*+\*\*++&\*+\*
  - <POST /checkout> \* + \* +& \*\*++&+
  - <DELETE /sessions> &\*+





**Cooperative VS Preemptive**

# Event-driven Programming

- Program flow determined by events
- Characterized by an event handler, main loop or reactor
- The reactor typically 'owns' the process
- `loop { button.on(:click).dispatch }`
- Binds event handlers to events
- Deep nesting of callbacks and constructs  
eg. Twisted for Python

# Batch programming

- wysiwig
- Flow determined by the developer
- Read → process → output

Event-driven  
execution,  
synchronous API.

# Database ...

```
class FiberedMysqlConnection
  def query( sql )
    send_query sql
    Fiber.yield
    get_result
  end
end
```

```
EM.run{ conn.query( 'SELECT SLEEP 1' ) }
```

... app server

```
class UsersController < AC::Base
  def index
    @users = User.paginate( :page => 1 )
  end
end
```

```
Thin::Server.start('0.0.0.0', 3000) do
  Fiber.new{ AC::Dispatcher.call( env ) }
```

```
end
```

# Frameworks

# Eventmachine

- Implementation of the Reactor pattern
- C++ with Ruby Bindings
- JRuby compatible through jeventmachine
- High performance event driven IO : epoll on Linux, kqueue on BSD and derivatives



# The Reactor

- Maintains a global current loop time
- Quantum is just under 100ms ... adjustable
- Verifies descriptors through a heartbeat
- Runs timer events
- New descriptors dropped onto a different queue
- Modified descriptors processed outside loop
- Stopped with special Loop Break Descriptor

# Bindings and Callbacks

- A portable signature registry / list, implemented as UUID identifiers
- Callback triggers :
  - timer fired
  - connection read
  - connection completed
  - connection accepted
  - connection unbound

# Eventable File Descriptors

- Nonblocking
- Notify of read and writability
- Closed in 3 ways
  - hard
  - immediate ( when possible, next tick )
  - after writing ( protocol handlers, serve HTTP page, close connection

# Deferrables

- A callback ... specifically a callback for :success or :failure
- Immediately pushed to the background and scheduled for future execution
- Deferrables should have an associated timeout

# Timers

- Periodic or oneshot
- Set a max timer count to avoid flooding the reactor
- Can be cancelled

# Files

- Very fast transfer of < 32k files with `EM#send_file_data`
- File watch API
  - Notified when modified, deleted or moved
- `fastfilereader` extension
  - Memory  $\leq$  disk transfer with `mmap`

# Processes

- Supports deferrable child processes through pipes
- EM#system
- Process watch API
  - Notified when forked or exits

# Client / Server

```
module Challenge
  def post_init
    send_data 'psw?'
  end

  def receive_data( data )
    ( @data ||= '' ) << data
  end
end

EM.run{ EM.connect 'localhost', 80, Challenge }
```



# Neverblock

# Why ?

- Most web apps block on IO
- Ruby processes has 'high' memory overheads
- 60MB RSS blocking 30ms on IO == \$\$\$
- Allocated in slabs – 80MB -> 120MB
- Significantly reduces capacity for concurrent requests
- eSpace stepped in

# How ?

- Fibered connection pool – typically 8 to 12
- Attaches to either EM or Rev as main loop
- Patches for Evented Mongrel && Thin to wrap requests in Fibers
- Fibered DB connections for postgres && mysql ( with mysqlplus )
- Respects transactions - BEGIN ... COMMIT required to be on the same connection




Connection 1 (idle)

Connection 2 (idle)

Fiber 1 is executing

Fiber 1  
(active  
)



Fiber 2  
(inactive)



Connection 1 (busy)

Connection 2 (idle)

Fiber 1  
(inactive)

SELECT ...

Fiber 2  
(active)


Fiber 1 issues a SQL query and passes control to Fiber 2



Connection 1 (busy)

Connection 2 (busy)

Fiber 1  
(inactive)



UPDATE ...

Fiber 2  
(inactive)

Fiber 2 issues a SQL query and both fibers are inactive



Connection 1 (idle)

Connection 2 (busy)

Fiber 1  
(active)

Alfred, 28, ...  
Tamer, 20, ...  
.  
.

Fiber 2  
(inactive)

Reactor notices first connection is finished so it resumes Fiber 1




**Connection 1 (idle)**

**Connection 2 (idle)**

**3600 rows affected**

**Fiber 1  
(done)**



**Fiber 2  
(active)**

**Fiber 1 is done and reactor notices second connection is finished  
so it resumes Fiber 2**



# Patches to Ruby IO

- Clever redirection to non-blocking methods
  - `IO#read*` => `IO#read_nonblock`
  - `IO#write*` => `IO#write_nonblock`
- Catches IO errors
  - `Errno::EWOULDBLOCK`
  - `Errno::EAGAIN`
  - `Errno::EINTR`
- Rescue attaches to the event loop
  - Requests notification of read / writability
  - Detach on success

# Operating System Events

# Signals

- Software interrupts
- Portable
  - `Signal.list` supported on all host systems
- Unobtrusive
  - Caught `Signal.trap( 'IO' ){ puts 'readable' }`
  - Sent with `Process.kill( 'IO', $ )`
- Can be scoped to a single process or a process group
- Daemons: `SIGHUP` && `SIGTERM`

# Handling Signals

- Ignore them – except KILL && STOP
- Catch with `trap('IO'){}`
- Let the default action apply
- SIGVTALRM is reserved for threading
- Signals the process every 10ms
- Flags that it's time to schedule
- Green threading maps to 1 OS thread, 1 per process, signals scoped per process

# Processes

- Process groups
  - Leader ? `Process#pid == process group identifier`
  - Group persist when leader exists as long as there's some members

# Multi-process Quality of Service

- Signals to enforce QOS
  - URG: Slow down
  - CONT: Resume normal operation
  - XCPU: Crawl
- Monitoring agent sends URG during load spike
- Sends CONT when load drops
- XCPU during extremely high load, effectively pausing the worker(s)

# POSIX Realtime Extensions

- Well supported, but implementations differ
  - Linux Kernel 2.6
  - FreeBSD / Darwin
  - Solaris
- Kernel level asynchronous IO
- Supports thread, signal based or no-op callbacks
- Reduces syscall overheads in IO bound applications

# How it works

- Control block
  - Struct: fd, buffer, notification, offset, number of bytes
- `aio_read( '/tmp/file' )` # signal notification with SIGIO
- Returns right away, faster than `O_NONBLOCK`
- Kernel signals with SIGIO on completion
  - Handler notification contains pointer to the the original control block



# Why it matters

- Parallel execution of `aio_read` && `aio_write`
  - `lio_listio( LIO_WRITE, list_of_cbs, 90, &list_sig )`
- A single user to kernel space switch
- 90 operations, a single signal notification
- Experimental `rb_aio` extension
- `AIO.write( { :filename => 'buffer' } )`

# Project Darkstar

<http://www.projectdarkstar.com>

# Gaming @ a Rails conf ?

- Rails Rack compatibility
- Middleware + ESI
- SOA architecture with thin processes + layers
- Russian doll pattern

# Architecture Overview

- Communications
  - Pubsub & direct channels supported
- Execution Kernel
  - Stateless execution of tasks
- Data storage
  - Transactional storage of serialized objects

# Task Execution

- Spawned and managed by a task manager
- Appears monothreaded, executes in parallel
- Immediate, delayed or periodic
- Parent  $\leq \geq$  child relationships
- Guaranteed to execute in order
- Children inherits priority of parent

# Task Lifetime

- Max execution time is 100ms - configurable
- Split long running tasks into multiple smaller tasks

# Managed Objects && References

- Managed object == a persitable entity
- References
  - Encapsulates all interactions with the data store
- Retrieval strategies
  - Get with the intention of modifying state
  - Get for read
  - Object#mark\_for\_update to promote a read only instance

# Listeners

- **Application listener**
  - Initialize: setup any game play items when first booted
  - Login callback: setup per player game state
- **Client listener**
  - Defines methods to repond to client events eg. logout



# Testing

# Test strategies

- Define environments : test / production
- Ability to switch between an evented and batched programming model depending on the environment
- A desirable side effect of loose coupling
- Very well suited for integration testing

# Block in testing mode

- Inject correlation middleware when in the test environment
- extend BlockInTest
- A single event may spawn \*data
  - @broker.collateral\_report( 'ACCXXXX' )
  - => [#<Report:0x202948>, #<Report:0x202234>]
- A single event may spawn many others
  - @broker.collateral\_report( 'ACCXXXX' )
  - => [:acknowledge\_reports, :reports]

# Event => expectation

- Mapping Event X to Reaction \*Y : Event IDXXXXXXXX => { IDXXXXXXXX => [Y,Y,Y] }
- Mapping Event X to Data Elements \*Z : Event IDXXXXXXXX => { IDXXXXXXXX => [Z,Z,Z] }
- Backbone of integration tests

# Existing specs ?

- [github.com/tmm1/em-spec](https://github.com/tmm1/em-spec)
- Bacon and Rspec backends
- A reactor / event loop for context
- How ?
  - Context executes in a Fiber, each expectation yields
- Unobtrusive
  - describe 'This is a context'
  - Em.spec 'This is an evented context'



# Questions ?

# Slides + Code

[http://github.com/methodmissing/railswaycon\\_events](http://github.com/methodmissing/railswaycon_events)

Thanks! Fork away.  
Follow  
[@methodmissing](#)  
on Twitter