

Lourens Naudé, Trade2Win

Embracing Events

Stop polling ... react.



RailsWay**Con**

Bio

- <http://github.com/methodmissing>
 - Ruby contractor
 - General misfit for the current economy
 - Ruby, C#, Java, admin 500 nodes, CSS * deferred payment
 - Currently building out a consumer forex platform @ Trade2Win
Slippage matter.
- creator
 - scrooge, mysqlplus_adapter, mri_instrumentation
- contributor
 - Rails, mysqlplus, query_memcached, thinking-sphinx, em-spec

The logo consists of the lowercase letters 't2w' in a bold, italicized, sans-serif font. The letters are red with a slight 3D effect, appearing to float above the background. The background is light gray with several thin, curved, white and light orange lines that suggest motion or a stylized railway track.

t2w

RailsWayCon

e-vent [i-vent]

something that happens or is regarded as happening; an occurrence, esp. one of some importance.

A change in state ...

- MySQL server sends results down the wire
 - Connection file descriptor becomes readable
 - Process results
 - Render results
- Ignorance is bliss
 - Never be flooded
 - Only handle events that matter

Talk#to_a

- Context switches
- Scheduling
- Programming models
- Frameworks
 - Eventmachine
 - Neverblock
- Operating System
- Lessons from massively multiplayer online games
- Testing
- Q & A ?

Getting Fibers (coroutines) and Threads

Scenario we all can relate to ...

- Quiet weekend day, erratic thoughts from the week before comes together
- Sharing space with your better half ... who's being chatty
- You're being interrupted at random
- Let's assume an average "in the zone" startup threshold of 2 minutes per interruption
- Deadlock's not an option ...

Multithreading

- Let the initial pattern of thought and the interruption be distinct Threads
- The interruption's effective right away
- The context switch is 'expensive' : coding VS whichever chore's been neglected
- Preemptive scheduling == constant interruption

Fiber / coroutine pool

- Two distinct Fibered contexts
- Able to defer the interruption
- Cheap switching : being thick skinned buys time
- You#resume the interruption if and when there's time to handle it

Thread

- Scheduling
 - Handled by the VM – timesliced @ 10ms by MRI
 - Executes as soon as CPU resources become available
- Availability
 - MRI 18: Maps to a single native thread
 - MRI 19: 1 x native thread for each Ruby thread - GIL
 - JRuby: 1 x Java thread for each Ruby thread
- Overheads
 - Large initial stack size
 - MRI uses a SIGVTALRM to signal a context switch

Fiber

- Scheduling
 - No concurrency – scheduled by application code
 - Never executes right away
- Availability
 - MRI 18: Patch by Aman Gupta & Joe Damato
 - MRI 19: Core Fiber class
 - JRuby: Poor Man's Fibers
- Overheads
 - Small initial stack size @ 4k
 - Context switches exclusively in user space – very fast

Why Fibers matter

- Computes a partial result - generator
- Yields back to the caller
- Saves execution state
- Caller resumes
- Data exchange
 - Accept arguments through `Fiber#resume`
 - Return values through `Fiber.yield`

Threads as Coroutines

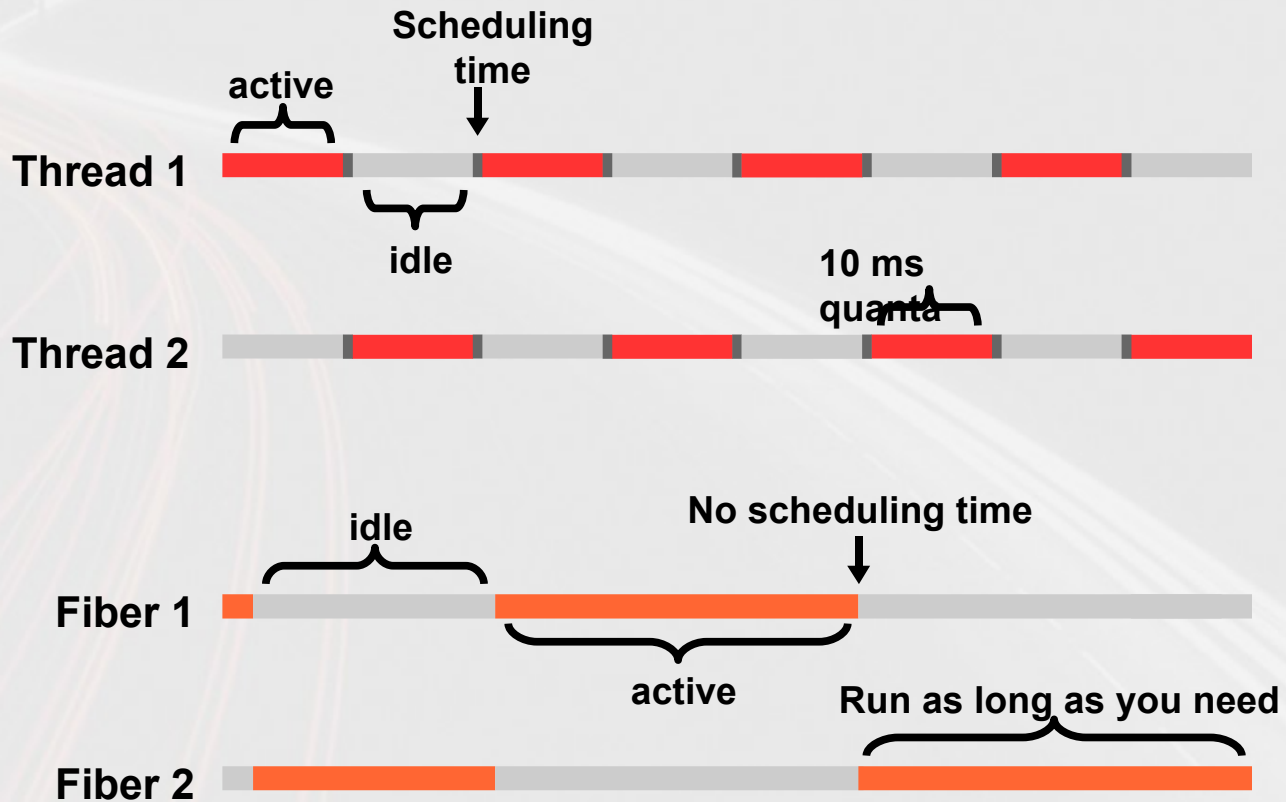
- Through a voodoo combination of
 - Thread priorities
 - Thread.pass
 - Thread.stop
 - Thread#run (yields CPU) or Thread#wakeup (no CPU)
- Poor Man's Fibers piggy backs off a Queue
- MRI 18 Fibers = Threads – preemptive scheduling, much like continuations, but with a smaller stack.
- Performant, but not as lightweight as Fibers

Cooperative scheduling

- Request / response with evented DB access
 - HTTP request wrapped in a Fiber
 - => mysql->fd : SELECT * FROM users WHERE id = 1
 - Execution paused
 - <= mysql->fd : #<User id: 1 uname: methodmissing>
 - Execution resumed
 - => mysql->fd: SELECT * FROM wishlist WHERE user_id = 1
 - Execution paused
 - <= mysql->fd: #<Wishlist id: 20 user_id 1>

Cooperative VS preemptive

- Fibered dispatch, query = *, result = +, int = &
 - <GET /orders> * + **++**++**++
 - <POST /checkout> * + * + **++
 - <DELETE /sessions> *+
- Threaded dispatch, interrupted every 10ms
 - <GET /orders> * & + **++&*+**++&*+*
 - <POST /checkout> * + * +& **++&+
 - <DELETE /sessions> &*+



Cooperative VS Preemptive

Event-driven Programming

- Program flow determined by events
- Characterized by an event handler, main loop or reactor
- The reactor typically 'owns' the process
- `loop { button.on(:click).dispatch }`
- Binds event handlers to events
- Deep nesting of callbacks and constructs
eg. Twisted for Python

Batch programming

- wysiwig
- Flow determined by the developer
- Read → process → output

Event-driven
execution,
synchronous API.

Database ...

```
class FiberedMysqlConnection
  def query( sql )
    send_query sql
    Fiber.yield
    get_result
  end
end
```

```
EM.run{ conn.query( 'SELECT SLEEP 1' ) }
```

... app server

```
class UsersController < AC::Base
  def index
    @users = User.paginate( :page => 1 )
  end
end
```

```
Thin::Server.start('0.0.0.0', 3000) do
  Fiber.new{ AC::Dispatcher.call( env ) }
```

```
end
```

Frameworks

Eventmachine

- Implementation of the Reactor pattern
- C++ with Ruby Bindings
- JRuby compatible through jeventmachine
- High performance event driven IO : epoll on Linux, kqueue on BSD and derivatives

The Reactor

- Maintains a global current loop time
- Quantum is just under 100ms ... adjustable
- Verifies descriptors through a heartbeat
- Runs timer events
- New descriptors dropped onto a different queue
- Modified descriptors processed outside loop
- Stopped with special Loop Break Descriptor

Bindings and Callbacks

- A portable signature registry / list, implemented as UUID identifiers
- Callback triggers :
 - timer fired
 - connection read
 - connection completed
 - connection accepted
 - connection unbound

Eventable File Descriptors

- Nonblocking
- Notify of read and writability
- Closed in 3 ways
 - hard
 - immediate (when possible, next tick)
 - after writing (protocol handlers, serve HTTP page, close connection

Deferrables

- A callback ... specifically a callback for :success or :failure
- Immediately pushed to the background and scheduled for future execution
- Deferrables should have an associated timeout

Timers

- Periodic or oneshot
- Set a max timer count to avoid flooding the reactor
- Can be cancelled

Files

- Very fast transfer of < 32k files with `EM#send_file_data`
- File watch API
 - Notified when modified, deleted or moved
- `fastfilereader` extension
 - Memory \leq disk transfer with `mmap`

Processes

- Supports deferrable child processes through pipes
- EM#system
- Process watch API
 - Notified when forked or exits

Client / Server

```
module Challenge
  def post_init
    send_data 'psw?'
  end

  def receive_data( data )
    ( @data ||= '' ) << data
  end
end

EM.run{ EM.connect 'localhost', 80, Challenge }
```


Observing the Reactor

(Interactive dtrace session)

Neverblock

Why ?

- Most web apps block on IO
- Ruby processes has 'high' memory overheads
- 60MB RSS blocking 30ms on IO == \$\$\$
- Allocated in slabs – 80MB -> 120MB
- Significantly reduces capacity for concurrent requests
- eSpace stepped in

How ?

- Fibered connection pool – typically 8 to 12
- Attaches to either EM or Rev as main loop
- Patches for Evented Mongrel && Thin to wrap requests in Fibers
- Fibered DB connections for postgres && mysql (with mysqlplus)
- Respects transactions - BEGIN ... COMMIT required to be on the same connection




Connection 1 (idle)

Connection 2 (idle)

Fiber 1 is executing

Fiber 1
(active
)



Fiber 2
(inactive)



Connection 1 (busy)

Connection 2 (idle)

Fiber 1
(inactive)

SELECT ...

Fiber 2
(active)


Fiber 1 issues a SQL query and passes control to Fiber 2



Connection 1 (busy)

Connection 2 (busy)

Fiber 1
(inactive)



UPDATE ...

Fiber 2
(inactive)

Fiber 2 issues a SQL query and both fibers are inactive



Connection 1 (idle)

Connection 2 (busy)

Fiber 1
(active)

Alfred, 28, ...
Tamer, 20, ...
.
.

Fiber 2
(inactive)

Reactor notices first connection is finished so it resumes Fiber 1




Connection 1 (idle)

Connection 2 (idle)

3600 rows affected

**Fiber 1
(done)**



**Fiber 2
(active)**

**Fiber 1 is done and reactor notices second connection is finished
so it resumes Fiber 2**

Patches to Ruby IO

- Clever redirection to non-blocking methods
 - `IO#read*` => `IO#read_nonblock`
 - `IO#write*` => `IO#write_nonblock`
- Catches IO errors
 - `Errno::EWOULDBLOCK`
 - `Errno::EAGAIN`
 - `Errno::EINTR`
- Rescue attaches to the event loop
 - Requests notification of read / writability
 - Detach on success

Operating System Events

Signals

- Software interrupts
- Portable
 - `Signal.list` supported on all host systems
- Unobtrusive
 - Caught `Signal.trap('IO'){ puts 'readable' }`
 - Sent with `Process.kill('IO', $)`
- Can be scoped to a single process or a process group
- Daemons: `SIGHUP` && `SIGTERM`

Handling Signals

- Ignore them – except KILL && STOP
- Catch with `trap('IO'){}`
- Let the default action apply
- SIGVTALRM is reserved for threading
- Signals the process every 10ms
- Flags that it's time to schedule
- Green threading maps to 1 OS thread, 1 per process, signals scoped per process

Processes

- Process groups
 - Leader ? `Process#pid == process group identifier`
 - Group persist when leader exists as long as there's some members

Multi-process Quality of Service

- Signals to enforce QOS
 - URG: Slow down
 - CONT: Resume normal operation
 - XCPU: Crawl
- Monitoring agent sends URG during load spike
- Sends CONT when load drops
- XCPU during extremely high load, effectively pausing the worker(s)

POSIX Realtime Extensions

- Well supported, but implementations differ
 - Linux Kernel 2.6
 - FreeBSD / Darwin
 - Solaris
- Kernel level asynchronous IO
- Supports thread, signal based or no-op callbacks
- Reduces syscall overheads in IO bound applications

How it works

- Control block
 - Struct: fd, buffer, notification, offset, number of bytes
- `aio_read('/tmp/file')` # signal notification with SIGIO
- Returns right away, faster than `O_NONBLOCK`
- Kernel signals with SIGIO on completion
 - Handler notification contains pointer to the the original control block

Why it matters

- Parallel execution of `aio_read` && `aio_write`
 - `lio_listio(LIO_WRITE, list_of_cbs, 90, &list_sig)`
- A single user to kernel space switch
- 90 operations, a single signal notification
- Experimental `rb_aio` extension
- `AIO.write({ :filename => 'buffer' })`

Project Darkstar

<http://www.projectdarkstar.com>

Gaming @ a Rails conf ?

- Rails Rack compatibility
- Middleware + ESI
- SOA architecture with thin processes + layers
- Russian doll pattern

Architecture Overview

- Communications
 - Pubsub & direct channels supported
- Execution Kernel
 - Stateless execution of tasks
- Data storage
 - Transactional storage of serialized objects

Task Execution

- Spawned and managed by a task manager
- Appears monothreaded, executes in parallel
- Immediate, delayed or periodic
- Parent $\leq \geq$ child relationships
- Guaranteed to execute in order
- Children inherits priority of parent

Task Lifetime

- Max execution time is 100ms - configurable
- Split long running tasks into multiple smaller tasks

Managed Objects && References

- Managed object == a persitable entity
- References
 - Encapsulates all interactions with the data store
- Retrieval strategies
 - Get with the intention of modifying state
 - Get for read
 - Object#mark_for_update to promote a read only instance

Listeners

- **Application listener**
 - Initialize: setup any game play items when first booted
 - Login callback: setup per player game state
- **Client listener**
 - Defines methods to repond to client events eg. logout

Testing

Test strategies

- Define environments : test / production
- Ability to switch between an evented and batched programming model depending on the environment
- A desirable side effect of loose coupling
- Very well suited for integration testing

Block in testing mode

- Inject correlation middleware when in the test environment
- extend BlockInTest
- A single event may spawn *data
 - @broker.collateral_report('ACCXXXX')
 - => [#<Report:0x202948>, #<Report:0x202234>]
- A single event may spawn many others
 - @broker.collateral_report('ACCXXXX')
 - => [:acknowledge_reports, :reports]

Event => expectation

- Mapping Event X to Reaction *Y : Event IDXXXXXXXX => { IDXXXXXXXX => [Y,Y,Y] }
- Mapping Event X to Data Elements *Z : Event IDXXXXXXXX => { IDXXXXXXXX => [Z,Z,Z] }
- Backbone of integration tests

Existing specs ?

- github.com/tmm1/em-spec
- Bacon and Rspec backends
- A reactor / event loop for context
- How ?
 - Context executes in a Fiber, each expectation yields
- Unobtrusive
 - describe 'This is a context'
 - Em.spec 'This is an evented context'



Questions ?

Slides + Code

http://github.com/methodmissing/railswaycon_events



Thanks! Fork away.